BroadMind: A Better Platforming Agent

Abstract

Recent work in reinforcement learning has focused on building generalist video game agents, as opposed to focusing on a particular genre of games. We aim to build a more specialized high-performance agent focused on the more challenging genre of platform games, which has received less attention. Utilizing symbolic representations of game state, we are training fully connected Neural Q-Network agents to successfully learn to play games with long term rewards and complex dynamics.

1. Background

Platform games, also known as platformers, are video games in which players guide a free-running avatar to jump between suspended platforms and avoid obstacles to advance through levels in the game. As a result of the array of environments to parse and the huge decision space, these games are very difficult for learning agents to play well. We are building an agent does not have to simulatenously recognize the screen-space pixels of the game and decide optimal policies. Instead, we have decoupled these two problems, and develop our algorithms in the Generalized Mario environment which provides symbolic state representations directly to the agent. However, even this symbolic representation is very large, motivating us to utilize neural network-based Q-learning approaches.

2. Approach

2.1. Experimental Setup

2.1.1. RL-GLUE

RL-Glue is a socket-based API that enables reinforcement learning experiments to work with software across multiple languages (Tanner & White, 2009). RL-Glue applications are made up of a core communications process, an agent, an environment, and an experiment. This allows our experiments to connect reusable Python agents across many open source environments written in languages including C++ and Java. It also enables us to customize our exper-

iments, such as optional game visualization, loading and saving trained policies, adjusting the game difficulty, etc.

2.1.2. GENERALIZED MARIO

Although our goal is to train agents to play Atari 2600 platforming games, it is easier to learn from a platform game that provides more layers of environment representations than raw screen pixels. To do this, we have started with the Generalized Mario environment that was contained in the 2009 AI Competition software package, and conveniently RL-Glue compatible (Togelius et al., 2010).

The Generalized Mario game has a total control input space of 12. This is made up of (-1, 0, 1) for left, none, and right motions respectively, on/off for jump, and on/off for run. We have used an encoding function that maps the integers 0 to 11 to and from this control input space.

Generalized Mario's interface provides many layers of state observation. It includes the (x,y) position and velocity of the mario actor, as well as a 22x16 tile grid semantically describing the screen space with the location of coins, pipes, blocks, etc. Separately, it has a list of all enemy positions on-screen, with similar position and velocity information as provided about Mario.

2.1.3. ARCADE LEARNING ENVIRONMENT

We have leveraged the Arcade Learning Environment, an open-source and RL-Glue compatible framework built on top of an Atari 2600 emulator (Bellemare et al., 2013). It provides an API for interfacing with the raw pixels, the current score, and the controller input of the game. This has allowed us to use original Atari games to train and evaluate our agents. We currently support 7 Atari platformers in our experimental setup:

1. Kung-Fu Master
2. Frostbite
3. Kangaroo
4. Pitfall!
5. Pitfall! 2
6. H.E.R.O.
7. Montezuma's Revenge

2.2. Learning Algorithms

2.2.1. STATE REPRESENTATION IN MARIO

It is challenging to find an effective representation of the Mario game state that enables effective Q-learning. The environment observations provided by Generalized Mario contains a wealth of symbolic information, but there are many possible encodings that we are investigating.

Currently, our representation is a tiled grid of integers, 20x12 tiles in size. The relative value in each tile is given as a "measure of goodness", enemies are -2, obstacles are -1, coins are +2, etc. The grid is centered on Mario's current location. We intend to represent the state with separate substrates for each class of objects, as in (Hausknecht et al., 2013), by breaking it out into separate background, enemy, reward, and actor layers. We hope to find that this representation will be universal across platform games.

It is worth mentioning that none of these representations encode the important velocity information about the actors, which is important in platformers where the characters move with some inertia. We may attempt an alternative approach that encodes this information.

2.2.2. NEURAL Q-LEARNING

Typically, Q-Learning approaches use a table representing the Q-function. This table contains a value for every combination of state and action. However, for very large state/action spaces such as platforming games, this is impractical as the space would take too many trials to explore and converge on an optimal policy. Even using optimizations such as nearest neighbor ran out of memory in our approach (spaces greater than 32 are not permitted, compared to at least 1,000+ for this application).

We have implemented a neural-network based Q-learning algorithm (see Algorithm 1) to allow us to learn on large state/action spaces with reasonable memory utilization by finding a useful hidden layer. Inspired by DeepMind's approach (Mnih et al., 2013), we have avoided multiple forward propagation steps when selecting optimal actions by using only the state as the input to the network, and a weight for each action as the output. Thus, we can select an optimal action for a state by propagating once, and selecting the max argument from the outputs. Also like DeepMind, we include an Experience Replay step that stores a history of events to re-learn. This avoids overfitting to the current situation and unlearning good behavior from earlier episodes. We are actively investigating methods to store experiences more intelligently. We can optionally use the standard Q-Learning update, or the SARSA calculation.

Algorithm 1 Neural Q-Network with Experience Replay

Initialize Neural Q-Network with random weights
Initialize Experience Pool to {}
for episode$= 1, m$ **do**
  Initialize previous state s_0 and previous action a_0 to NULL
  for $t = 1, T$ **do**
    Observe state s_t from the emulator
    With probability ϵ_a, select random action a_t
    Else, set $a_i = \max_a Q(s_t, a)$ by forward propagating s_t through Q
    if s_{t-1} and a_{t-1} are not NULL **then**
      Observe reward r from emulator
      With probability ϵ_r, store the experience $\{s_{t-1}, a_{t-1}, r, s_t, a_t\}$ in the pool
    end if
    for re-experience$= 1, ex$ **do**
      Randomly sample experience $\{s'_0, a'_0, r', s'_1, a'_1\}$ from the pool
      if using SARSA update rule **then**
        Compute $v = r' + \gamma Q(s'_1, a'_1)$
      else
        Compute $v = r' + \gamma \max_a Q(s'_1, a)$
      end if
      Update Q through backpropagation of value v on output a'_0 with state s'_0
    end for
    Apply action a_i to emulator
    Update state $s_{t-1} = s_t$ and action $a_{t-1} = a_t$
  end for
end for

2.2.3. EXTENSION TO ATARI PLATFORMERS

We have setup the ALE environment, and connected default agents to it. However, we have yet to attempt to port our Mario agents to these problems. Initially, we will use 3 colored pixel substrates as the state representation, as in (Hausknecht et al., 2013), but we hope to reconstruct an object layer as well.

2.3. Results

We have found that the need to assign initially random weights to the neural network can make the initial policies very flawed. To counter this, we use heavy exploration bias (ϵ 1.0) for early episodes, while transitioning to exploitation policies (ϵ 0.1) at later episodes. We have also biased the random action to prefer motion to the right, helping the agent to explore more of the game.

We trained agents over 1000 episodes in the Mario environment using the state encoded in the format described in section 2.2.1, and using the Neural Q-Network with expe-

Figure 1. Mario agent trained with a neural q-network with a hidden layer of 126 nodes. The agent was trained for 1000 episodes of the same level seed 3 and difficulty 1. The initial exploration factor was 1.0, and this decreased by 0.05 every 100 episodes, until stopping at 0.1. The total reward gained by the agent was summed over each episode, and the running average of 100 episodes is shown here.

rience replay of Algorithm 1. We use Mario level of difficulty 1 so that the levels provide a challenge with enemies, but not so hard that a random agent cannot make considerable progress.

We have yet to optimize our agents over the space of configuration parameters (learning rate, regularization, max number of experiences, probability of remembering an experience, discount factor on future rewards). Also, we have not yet attempted to train an agent on a particular level, and evaluate its performance on a new level of a different difficulty level.

Acknowledgments

References

Bellemare, Marc G, Naddaf, Yavar, Veness, Joel, and Bowling, Michael. The arcade learning environment: An evaluation platform for general agents. *Journal of Artificial Intelligence Research*, 47:253–279, 2013.

Hausknecht, M., Lehman, J., Miikkulainen, R., and Stone, P. A neuroevolution approach to general atari game playing. *IEEE Transactions on Computational Intelligence and AI in Games*, 2013.

Mnih, Volodymyr, Kavukcuoglu, Koray, Silver, David, Graves, Alex, Antonoglou, Ioannis, Wierstra, Daan, and Riedmiller, Martin. Playing atari with deep reinforcement learning. *NIPS Workshop on Deep Learning and Unsupervised Feature Learning*, 2013.

Tanner, Brian and White, Adam. Rl-glue: Language-independent software for reinforcement-learning experiments. *Journal of Machine Learning Research*, 10: 2133–2136, 2009.

Togelius, J., Karakovskiy, S., and Baumgarten, R. The 2009 mario ai competition. *Evolutionary Computation (CEC), 2010 IEEE Congress on.*, pp. 1–8, 2010.