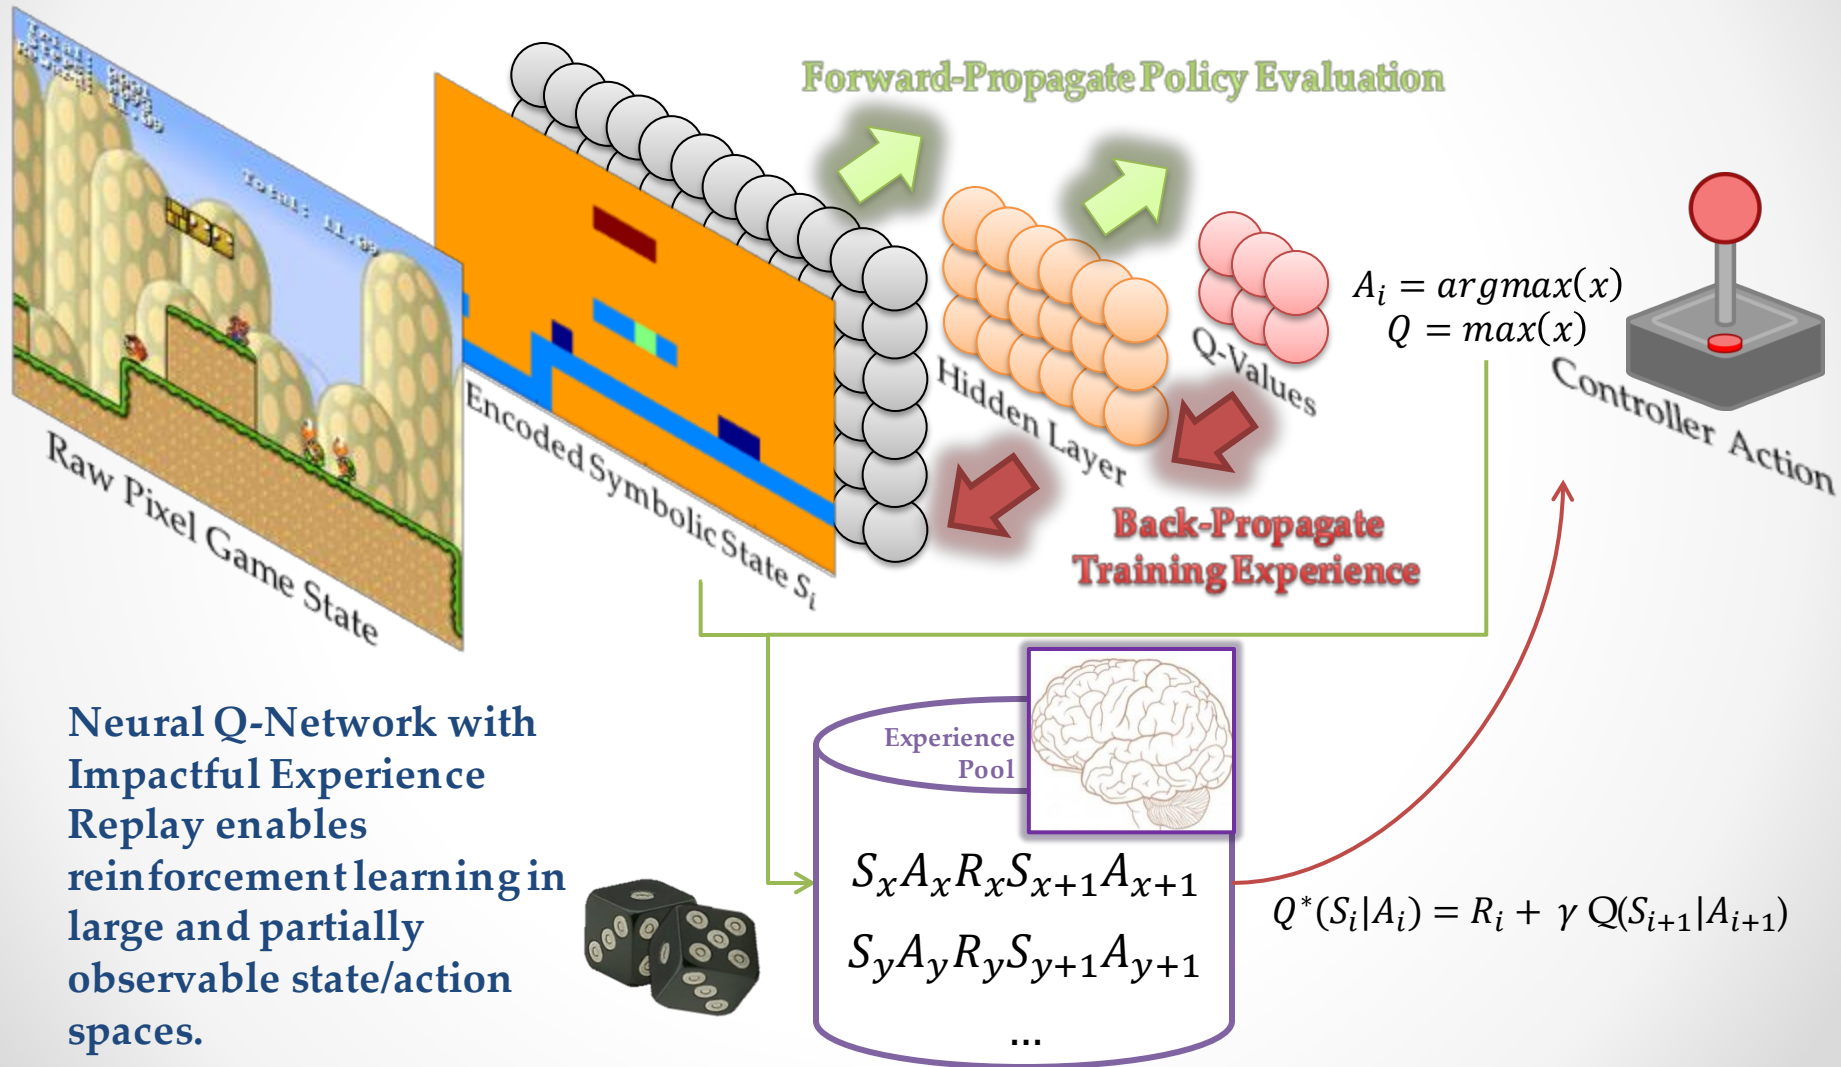
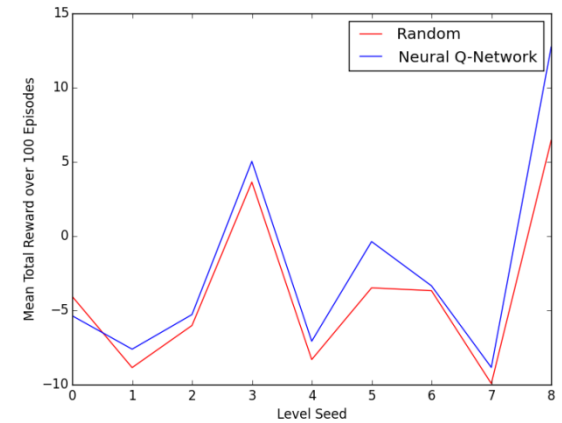




BroadMind - Approach



BroadMind - Results



Impact-weighted experience remembrance slows the learning rate, and achieves the same long-term performance when trained on only one Generalized Mario level.

Using a multiple-substrate state representation of background, enemy, and reward layers does not exceed random performance the way that a larger single layer can.

Agents trained on one Mario level achieve good performance when transferred to a new, unseen level.